


Route mechanics and evidence boundaries

run m2-one-shot

M2 compares baseline target emission with route outputs over identical fixture groups.



Current R3 result

The run shows a target-syntax lift (3/9 -> 9/9) but no admitted lift: baseline admitted accuracy is 0/9 and single_ir admitted accuracy is 0/9. Candidate verdict accuracy for single_ir is 3/9 and remains rejected evidence.