


Route mechanics and evidence boundaries

run m2-one-shot

M2 compares direct SMT emission with a bounded IR hop over identical fixture groups.




Current R3 result

The run shows a target-syntax lift (3/9 -> 9/9) but no admitted lift: both routes have 0/9 admission and 0/9 admitted verdict accuracy.
Candidate verdict accuracy for single_ir is 3/9 and remains rejected evidence.

M2 route metrics


Exact-ratio measurements over three groups and three seeds per route.

n = 9 rows per route




Failure taxonomy by route

Diagnostic categories are non-exclusive row hits over nine route rows.



Counts come from route_evaluation.route_category_counts. Categories can co-occur in one row, so columns are diagnostic burden rather than a partition of samples.



Per-source candidate coverage				
source	spans	rules	rejected	raw cache
real.jstage.thyroid_tumor_2024.cq_recommendations	4	4	0	complete
real.jstage.sepsis_dic_2025.guideline_article	3	0	6	complete