


Route mechanics and evidence boundaries

run m2-one-shot

M2 compares baseline target emission with route outputs over identical fixture groups.




Current R3 result

The run shows a target-syntax lift (3/9 -> 9/9) but no admitted lift: baseline admitted accuracy is 0/9 and single_ir admitted accuracy is 0/9. Candidate verdict accuracy for single_ir is 3/9 and remains rejected evidence.

M2 route matrix metrics


Exact-ratio measurements over identical groups and seeds per route.

n = 9 rows per route



Failure taxonomy by route

Diagnostic categories are non-exclusive row hits over each route's rows.




Counts come from route_evaluation.route_category_counts. Categories can co-occur in one row, so columns are diagnostic burden rather than a partition of samples.

Real guideline intake scope

no clinical claim

Public-source rows are candidate-only artifacts outside locked M1/M2 scoring.



Per-source candidate coverage				
source	spans	rules	rejected	raw cache
real.jstage.thyroid_tumor_2024.cq_recommendations	4	4	0	complete
real.jstage.sepsis_dic_2025.guideline_article	3	0	6	complete