


Route mechanics and evidence boundaries

run m3-routes

The route matrix compares baseline target emission with route outputs over identical fixture groups.




Route matrix result

The route matrix keeps route.direct_smt as baseline. Highest target syntax is ckc_layered at 4/4; candidate accuracy remains rejected-output audit evidence.

Route matrix metrics


Exact-ratio measurements over identical groups and seeds per route.

n = 4 rows per route



Failure taxonomy by route

Diagnostic categories are non-exclusive row hits over each route's rows.




Counts come from route_evaluation.route_category_counts. Categories can co-occur in one row, so columns are diagnostic burden rather than a partition of samples.

Real guideline intake scope

no clinical claim

Public-source rows are candidate-only artifacts outside locked M1/M2 scoring.




Per-source candidate coverage				
source	spans	rules	rejected	raw cache
real.jstage.thyroid_tumor_2024.cq_recommendations	4	4	0	complete
real.jstage.sepsis_dic_2025.guideline_article	3	0	6	complete

Deterministic pipeline comparison

run m3-compare; n = 8 groups

Layered CKC artifacts are compared with direct rule-to-SMT over identical M3 groups.



direct_rule

layered_ckc

Interpretation

Direct baseline verdict accuracy is 8/8. Layered verdict accuracy is 8/8; layered-minus-direct verdict delta is 0/1.

Component reuse delta is 51/59. This is deterministic pipeline evidence and is separate from model-route lift.