


Route mechanics and evidence boundaries

run m3-routes

The route matrix compares baseline target emission with route outputs over identical fixture groups.




Route matrix result

The route matrix keeps route.direct_smt as baseline. Highest target syntax is ckc_layered at 4/4; candidate accuracy remains rejected-output audit evidence.

Route matrix metrics


Exact-ratio measurements over identical groups and seeds per route.

n = 4 rows per route




Failure taxonomy by route

Diagnostic categories are non-exclusive row hits over each route's rows.



Counts come from route_evaluation.route_category_counts. Categories can co-occur in one row, so route profiles are diagnostic burden rather than a partition of samples.



Per-source candidate coverage				
source	spans	rules	rejected	raw cache
real.jstage.thyroid_tumor_2024.cq_recommendations	4	4	0	complete
real.jstage.sepsis_dic_2025.guideline_article	3	0	6	complete

Deterministic pipeline comparison

Layered CKC artifacts are compared with direct rule-to-SMT over identical M3 groups.

run m3-compare; n = 10 groups



direct_rule

layered_ckc

Interpretation



Direct baseline verdict accuracy is 10/10. Layered verdict accuracy is 10/10; layered-minus-direct verdict delta is 0/1.

Component reuse delta is 60/71. This is deterministic pipeline evidence and is separate from model-route baseline deltas.

Model-free coverage

run m3-coverage

Build-set mappings are applied to fresh fixture groups with runtime AI disabled.



Apply rows

- m2_holdout_conflict 12/12 components
- m3_metamorphic_conflict 12/12 components
- m3_exception_resolver 12/12 components

The model-per-document baseline would make 5 apply calls; this run makes zero apply-phase model calls.